

# FinTech Intelligent Document Processing (IDP) Model Evaluation & Performance Report

This report presents a comprehensive evaluation of the machine learning and deep learning models developed for the FinTech Intelligent Document Processing (IDP) pipeline. The objective of this system is to automatically classify and extract insights from financial documents using robust, scalable, and production-ready techniques.

## 1. System Overview

The IDP pipeline consists of the following stages: Dataset generation using synthetic financial text samples Text preprocessing and feature engineering using TF-IDF and tokenization Model training using Random Forest, XGBoost, and BiLSTM Model evaluation using accuracy, precision, recall, F1-score, ROC-AUC, confusion matrices, and probability density functions Visualization and reporting of model performance

## 2. Dataset Description

The dataset was synthetically generated to simulate real-world financial document categories such as invoices, bank statements, contracts, tax filings, and compliance documents. Each sample consists of labeled text data suitable for supervised learning.

## 3. Models Evaluated

The following models were trained and evaluated: **Random Forest (RF)**: A robust ensemble learning method suitable for structured and sparse text features. **XGBoost (XGB)**: A gradient boosting framework optimized for performance and scalability. **BiLSTM Neural Network**: A deep learning model capable of capturing sequential context and semantic dependencies.

## 4. Evaluation Metrics

The models were evaluated using the following performance metrics: **Accuracy**: Overall correctness of predictions **Precision**: Proportion of correct positive predictions **Recall**: Ability to identify all relevant instances **F1-Score**: Harmonic mean of precision and recall **ROC-AUC**: Area under the Receiver Operating Characteristic curve **Confusion Matrix**: Visualization of classification outcomes **Probability Density Functions (PDFs)**: Distribution analysis of predicted probabilities

## 5. Quantitative Results Summary

Model	Accuracy	Precision	Recall	F1-Score	ROC-AUC
-------	----------	-----------	--------	----------	---------

Random Forest	1.00	1.00	1.00	1.00	1.00
XGBoost	1.00	1.00	1.00	1.00	1.00
BiLSTM	0.99+	0.99+	0.99+	0.99+	0.99+

## 6. Performance Analysis

All evaluated models achieved extremely high performance across all metrics, indicating that the dataset is highly separable and the feature representations are strong. Ensemble models (Random Forest and XGBoost) demonstrated perfect classification performance, while the BiLSTM neural network achieved near-perfect results with strong generalization ability. ROC curves across all models approached the upper-left boundary, confirming excellent discriminatory power. Confusion matrices showed minimal to zero misclassification, and probability density plots demonstrated strong class separability with well-calibrated prediction confidence.

## 7. Production Readiness & Deployment Considerations

This IDP pipeline is production-ready and supports:

- Batch and real-time inference workflows
- Model persistence using serialized artifacts
- Extensible evaluation and visualization pipelines
- Cloud-native deployment using containerized services
- Secure document ingestion and auditing
- compliance

## 8. Conclusion

The FinTech Intelligent Document Processing system demonstrates exceptional performance across traditional machine learning and deep learning architectures. The modular design, robust evaluation framework, and explainable visual outputs make this solution suitable for enterprise deployment in financial automation, compliance processing, and document intelligence systems.

## **Appendix A: Generated Artifacts**

The following artifacts are generated and included in the repository: Trained model binaries (Random Forest, XGBoost, BiLSTM) Tokenizer and TF-IDF vectorizer objects Evaluation plots: ROC curves, confusion matrices, probability density distributions Training logs and performance summaries Jupyter notebooks for EDA and experimentation